

DD2448 Foundations of Cryptography

Lecture 5

Douglas Wikström
KTH Royal Institute of Technology
dog@kth.se

March 25, 2020

Second Pre-Image Resistance

Definition. A function $h : \{0, 1\}^* \rightarrow \{0, 1\}^*$ is said to be **second pre-image resistant** if for every polynomial time algorithm A and a random x

$$\Pr[A(x) = x' \wedge x' \neq x \wedge f(x') = f(x)] < \epsilon(n)$$

for a negligible function ϵ .

Note that A is given not only the output of f , but also the **input** x , but it must find a **second** pre-image.

Definition. Let $f = \{f_\alpha\}_\alpha$ be an ensemble of functions. The “function” f is said to be **collision resistant** if for every polynomial time algorithm A and randomly chosen α

$$\Pr[A(\alpha) = (x, x') \wedge x \neq x' \wedge f_\alpha(x') = f_\alpha(x)] < \epsilon(n)$$

for a negligible function ϵ .

Definition. Let $f = \{f_\alpha\}_\alpha$ be an ensemble of functions. The “function” f is said to be **collision resistant** if for every polynomial time algorithm A and randomly chosen α

$$\Pr[A(\alpha) = (x, x') \wedge x \neq x' \wedge f_\alpha(x') = f_\alpha(x)] < \epsilon(n)$$

for a negligible function ϵ .

An algorithm that gets a small “advice string” for each security parameter can easily hardcode a collision for a fixed function f , which explains the random index α .

Relations for Compressing Hash Functions

- ▶ If a function is not second pre-image resistant, then it is not collision-resistant.

Relations for Compressing Hash Functions

- ▶ If a function is not second pre-image resistant, then it is not collision-resistant.
 1. Pick random x .
 2. Request second pre-image $x' \neq x$ with $f(x') = f(x)$.
 3. Output x' and x .

Relations for Compressing Hash Functions

- ▶ If a function is not second pre-image resistant, then it is not collision-resistant.
 1. Pick random x .
 2. Request second pre-image $x' \neq x$ with $f(x') = f(x)$.
 3. Output x' and x .
- ▶ If a function is not one-way, then it is not second pre-image resistant.

Relations for Compressing Hash Functions

- ▶ If a function is not second pre-image resistant, then it is not collision-resistant.
 1. Pick random x .
 2. Request second pre-image $x' \neq x$ with $f(x') = f(x)$.
 3. Output x' and x .
- ▶ If a function is not one-way, then it is not second pre-image resistant.
 1. Given random x , compute $y = f(x)$.
 2. Request pre-image x' of y .
 3. Repeat until $x' \neq x$, and output x' .

Random Oracles

Random Oracle As Hash Function

A random oracle is simply a randomly chosen function with appropriate domain and range.

A random oracle is the **perfect** hash function. Every input is mapped **independently** and **uniformly** in the range.

Let us consider how a random oracle behaves with respect to our notions of security of hash functions.

Pre-Image of Random Oracle

We assume with little loss that an adversary always “knows” if it has found a pre-image, i.e., it queries the random oracle on its output.

Theorem. Let $H : X \rightarrow Y$ be a randomly chosen function and let $x \in X$ be randomly chosen. Then for every algorithm A making q oracle queries

$$\Pr[A^{H(\cdot)}(H(x)) = x' \wedge H(x) = H(x')] \leq 1 - \left(1 - \frac{1}{|Y|}\right)^q.$$

Pre-Image of Random Oracle

We assume with little loss that an adversary always “knows” if it has found a pre-image, i.e., it queries the random oracle on its output.

Theorem. Let $H : X \rightarrow Y$ be a randomly chosen function and let $x \in X$ be randomly chosen. Then for every algorithm A making q oracle queries

$$\Pr[A^{H(\cdot)}(H(x)) = x' \wedge H(x) = H(x')] \leq 1 - \left(1 - \frac{1}{|Y|}\right)^q.$$

Proof. Each query x' satisfies $H(x') \neq H(x)$ independently with probability $1 - \frac{1}{|Y|}$.

Second Pre-Image of Random Oracle

We assume with little loss that an adversary always “knows” if it has found a second pre-image, i.e., it queries the random oracle on the input and its output.

Theorem. Let $H : X \rightarrow Y$ be a randomly chosen function and let $x \in X$ be randomly chosen. Then for every such algorithm A making q oracle queries

$$\Pr[A^{H(\cdot)}(x) = x' \wedge x \neq x' \wedge H(x) = H(x')] \leq 1 - \left(1 - \frac{1}{|Y|}\right)^{q-1}.$$

Second Pre-Image of Random Oracle

We assume with little loss that an adversary always “knows” if it has found a second pre-image, i.e., it queries the random oracle on the input and its output.

Theorem. Let $H : X \rightarrow Y$ be a randomly chosen function and let $x \in X$ be randomly chosen. Then for every such algorithm A making q oracle queries

$$\Pr[A^{H(\cdot)}(x) = x' \wedge x \neq x' \wedge H(x) = H(x')] \leq 1 - \left(1 - \frac{1}{|Y|}\right)^{q-1}.$$

Proof. Same as pre-image case, except we must waste one query on the input value to get the target in Y .

Collision Resistance of Random Oracles

We assume with little loss that an adversary always “knows” if it has found a collision, i.e., it queries the random oracle on its outputs.

Theorem. Let $H : X \rightarrow Y$ be a randomly chosen function. Then for every such algorithm A making q oracle queries

$$\begin{aligned}\Pr[A^{H(\cdot)} = (x, x') \wedge x \neq x' \wedge H(x) = H(x')] &\leq 1 - \prod_{i=1}^{q-1} \left(1 - \frac{i}{|Y|}\right) \\ &\leq \frac{q(q-1)}{2|Y|} .\end{aligned}$$

Collision Resistance of Random Oracles

We assume with little loss that an adversary always “knows” if it has found a collision, i.e., it queries the random oracle on its outputs.

Theorem. Let $H : X \rightarrow Y$ be a randomly chosen function. Then for every such algorithm A making q oracle queries

$$\begin{aligned}\Pr[A^{H(\cdot)} = (x, x') \wedge x \neq x' \wedge H(x) = H(x')] &\leq 1 - \prod_{i=1}^{q-1} \left(1 - \frac{i}{|Y|}\right) \\ &\leq \frac{q(q-1)}{2|Y|} .\end{aligned}$$

Proof. $1 - \frac{i-1}{|Y|}$ bounds the probability that the i th query does not give a collision for any of the $i - 1$ previous queries, conditioned on no previous collision.

Iterated Hash Functions

Suppose that we are given a collision resistant hash function

$$f : \{0,1\}^{n+t} \rightarrow \{0,1\}^n .$$

How can we construct a collision resistant hash function

$$h : \{0,1\}^* \rightarrow \{0,1\}^n$$

mapping any length inputs?

Construction.

1. Let $x = (x_1, \dots, x_k)$ with $|x_i| = t$ and $0 < |x_k| \leq t$.
2. Let x_{k+1} be the total number of bits in x .
3. Pad x_k with zeros until it has length t .
4. $y_0 = 0^n$, $y_i = f(y_{i-1}, x_i)$ for $i = 1, \dots, k + 1$.
5. Output y_{k+1}

Here the total number of bits is bounded by $2^t - 1$, but this can be relaxed.